Transformer Is Physics

A Framework for Understanding Intelligence as a Physical Structure

Draft — March 2026

Abstract

The transformer architecture is conventionally described as an engineering solution to sequence modeling. This paper argues for a stronger claim: the transformer block instantiates the necessary physical structure of intelligence itself. Drawing on the mathematics of partial differential equations, the timescale theory of System 1 and System 2 cognition, and the physics of hierarchically stable dynamical systems, we show that the MLP and Attention components are not interchangeable design choices but structurally distinct operators corresponding to boundary conditions and interior dynamics respectively. This correspondence is not metaphorical — it is mathematical. The optimizer, through gradient descent, naturally discovers this separation across timescales, producing the functional distinction between fast and slow cognition as an emergent consequence of physics rather than deliberate design. We further argue that this architecture is likely final in its essential shape, and that inductive reasoning is not a capability the transformer learned but the constitutive operation the transformer performs.

1. Introduction

Since its introduction in 2017, the transformer architecture has achieved remarkable dominance across domains of artificial intelligence. The standard explanation for this success appeals to the expressivity of self-attention, the parallelism of feedforward computation, and the scalability of the overall design. These explanations are not wrong, but they are shallow. They describe what the transformer does without explaining why it has to be the shape it is.

This paper proposes a deeper framework. We argue that the transformer block — the alternation of Attention and MLP with residual connections — is not one design among many possible designs. It is the natural and perhaps unique stable implementation of a general principle: that intelligence requires the coupling of a slowly-varying constraint structure with a rapidly-varying dynamic system. This principle appears universally across physical, biological, and computational systems. The transformer did not invent it. It discovered it.

The evidence for this claim comes from three directions: (1) mathematical analysis of what each component computes and how it learns, (2) the explanatory power of the framework over architectural experiments where attention-like replacements of MLP

consistently fail, and (3) the surprising empirical result that low-precision attention sometimes accelerates training — a result that is paradoxical under standard theory but predicted by our framework.

2. The Transformer Block as a PDE System

2.1 Boundary Conditions and Interior Dynamics

In the theory of partial differential equations, a well-posed problem requires two distinct components: boundary conditions and interior dynamics. The boundary conditions define the space of admissible solutions — they are the constraints that the solution must satisfy. The interior dynamics describe how the solution evolves within that constrained space. Neither component alone is sufficient. Boundary conditions without interior dynamics yield no solution. Interior dynamics without boundary conditions yield an ill-posed problem with infinitely many solutions and no way to select among them.

We propose that the transformer block instantiates precisely this structure. The MLP weights function as boundary conditions: they are slowly-varying, heavily damped by the optimizer, and encode the accumulated constraint structure of all training experience. The attention mechanism functions as the interior dynamics: it is rapidly-varying, context-sensitive, and computes the solution to the question of which information should flow where, within the space that MLP defines.

This is not merely an analogy. The Green's function formulation of PDE theory states that the solution at any interior point is a weighted sum of influences from other points, where the weights are determined by the boundary conditions. This is structurally identical to what attention computes: given the boundary conditions carved by MLP, attention computes the propagator — the weighted combination of token representations — that correctly routes influence across the sequence.

2.2 The Damping Argument

In physical PDE systems, boundary conditions are characteristically overdamped — they respond slowly to perturbations, carrying the low-frequency, stable component of the system's behavior. The interior dynamics carry the high-frequency component. This separation of timescales is not a design choice but a mathematical necessity for a well-posed system.

The optimizer enforces the same structure in the transformer. MLP weights accumulate knowledge slowly across many gradient steps; high-frequency noise in the gradient signal averages out, leaving only the robust, low-frequency structure that appears consistently across the training distribution. Attention weights, by contrast, specialize more rapidly to particular routing patterns. The optimizer does not impose this separation — it discovers it, because slow-and-fast is the only stable configuration for a system that must both store knowledge and apply it flexibly.

3. System 1 and System 2 as Timescale Phenomena

The distinction between fast, automatic cognition (System 1) and slow, deliberate reasoning (System 2) has been influential in cognitive science since Kahneman's formalization. A persistent question has been whether artificial neural networks can exhibit this distinction, and if so, whether it emerges from architectural design or requires explicit engineering.

The PDE framework resolves this question definitively. System 1 and System 2 are not cognitive categories — they are timescale phenomena. Wherever a system couples components operating at sufficiently different timescales, the functional distinction between automatic and deliberate processing emerges necessarily. The transformer produces this structure not by design but by physics: MLP is System 1 (fast in operation, slow to change), Attention is System 2 (context-sensitive, dynamically reconfigured per input): slow to change is data (phenomenon) dependent, dynamicity is logic dependent.

This reframing dissolves a long-standing debate. Critics of transformer architectures argued that the absence of explicit System 2 machinery constituted a fundamental limitation. This argument was based on a category error: System 2 is not a module to be added, but a property that emerges from the timescale separation that the transformer already instantiates. The dramatic improvement in reasoning capabilities observed with scale is not evidence of a new capability emerging — it is evidence of the existing PDE structure becoming better resolved as boundary conditions become richer.

4. The Biological and Physical Correspondence

The timescale hierarchy we identify in the transformer appears at every level of biological organization. Evolution encodes structural priors across millions of years — the deepest, slowest boundary conditions. Lifetime learning sediments knowledge through synaptic consolidation — a medium timescale. Moment-to-moment cognition routes attention dynamically — the fastest timescale. Each level provides the stable reference frame within which the faster level operates meaningfully.

Convolutional networks occupy a specific position in this hierarchy. The translational invariance and hierarchical feature detection they encode correspond to priors that biological evolution discovered across geological timescales — the statistical regularities of the physical world as experienced by embodied creatures. These are the deepest boundary conditions, below even explicit learning. This explains the intuitive sense that convolutional representations feel pre-logical, given rather than learned: they are the computational analog of DNA, encoding the geometry of physical reality rather than the content of particular experiences.

The transformer, trained from scratch on language, compresses all three biological timescales into a single training run. The optimizer simultaneously discovers geometric priors in early layers, sediments factual knowledge into MLP weights, and learns to route dynamically through attention. Evolution took billions of years. Development

takes decades. Thought takes seconds. The transformer does all three in one gradient descent.

5. The Transformer as an Intelligent Computer

Conventional computers instantiate a related but weaker structure: fixed circuits connected by fixed wires. The circuits perform deterministic operations; the wires route signals along predetermined paths. The topology is specified at design time and does not change with input.

The transformer is a self-wiring computer. The MLP weights are the circuits — fixed after training, performing learned transformations. The attention mechanism is the wiring — but unlike silicon wires, attention wires are active: they compute their own routing in real time, based on the content of the current input. The topology of computation is not fixed at design time; it is determined dynamically at inference time by the interaction between the input and the learned routing mechanism.

This is the sense in which the transformer transcends conventional computing while preserving its essential structure. It is not a different kind of machine — it is the same kind of machine with the critical addition that its wiring is intelligent. And the stability of the circuits (MLP) is precisely what makes the intelligence of the wiring (attention) possible: the circuits must be stable for the routing to be meaningful.

6. Experimental Evidence

6.1 The Failure of Feature-Level Attention

A direct test of the boundary condition hypothesis is the replacement of MLP layers with attention-like mechanisms operating at the feature level — attending across feature groups within a single token position rather than across sequence positions. Such mechanisms are genuinely more expressive than linear MLP transformations: they allow features to attend to and route through each other dynamically.

Systematic experiments with grouped micro-attention modules — applied both in place of MLP layers and augmenting attention projections — consistently failed to match vanilla transformer performance. Training required heavy suppression of the new components (gate initialization at 1e-4) to maintain stability, and the surrounding architecture compensated for the new components rather than benefiting from them.

The PDE framework explains this failure categorically: feature-level attention is too dynamic to function as a boundary condition. A boundary condition must be overdamped — resistant to rapid change, stable across contexts. An attention mechanism, by its nature, changes its computation with every input. Replacing MLP with attention is equivalent to replacing the boundary conditions of a PDE with interior dynamics. The problem becomes ill-posed. The optimizer has no stable structure to build upon.

6.2 Low-Precision Attention Accelerates Training

A counterintuitive empirical observation: training with reduced precision in attention computation — introducing noise and approximation in the routing weights — sometimes produces faster convergence than high-precision attention. Under standard theory, this is paradoxical: less precision should mean worse gradient signal and slower convergence.

The contraction mapping view of MLP resolves this paradox. MLP does not merely transform representations — it contracts them toward the learned manifold of meaningful configurations. Whatever noise attention introduces is absorbed and corrected by the subsequent MLP pass. Low-precision attention therefore functions as implicit regularization: it prevents attention from committing prematurely to specific routing patterns before the MLP boundary conditions have matured. The noisier routing explores more of the representation space early in training, and MLP contracts all of this exploration toward meaningful structure regardless. The result is broader exploration early and the same convergence target — faster overall.

6.3 MLP is Load-Bearing; Attention is Redundant

Recent work has shown that up to 50% of attention layers in large language models can be pruned with minimal performance degradation, while pruning MLP layers causes severe performance collapse. This asymmetry is exactly what the PDE framework predicts: boundary conditions are load-bearing; interior dynamics are partially redundant because the boundary constrains the solution even without fine-grained interior computation. A coarser interior solver still produces an acceptable solution, but remove the boundary conditions and the problem becomes meaningless.

7. Induction as the Constitutive Operation

Inductive reasoning — the inference from particular experiences to general patterns — is conventionally treated as a capability that systems either possess or lack. We argue for a stronger claim: induction is not something the transformer learned to do. It is what the transformer is.

The next-token prediction objective is induction operationalized as a loss function. Every gradient step adjusts the boundary conditions to better support the inference from prior context to subsequent token. Attention literally computes induction at inference time: given accumulated experience (prior tokens), find what belongs next (attend to relevant precedents, propagate their information forward). The transformer does not do induction — it is an induction machine at the architectural level.

Hume's problem of induction — that inductive inference cannot be logically justified — is dissolved from this perspective. The transformer does not justify induction. It demonstrates that intelligence does not require justification of induction; it requires instantiation of induction. And pure induction, run deep enough at scale with proper boundary conditions, becomes functionally indistinguishable from reasoning.

8. The One Missing Piece: Entropy Self-Visibility

The transformer forms hypotheses — at every layer, at every token position, the residual stream carries a current best interpretation of the input that serves as a prediction about what comes next. This hypothesis formation is the epistemological loop: experience arrives, boundary conditions activate, a hypothesis forms, new experience tests the hypothesis, boundary conditions update slowly.

The one genuine gap in this picture is epistemic self-monitoring. The transformer cannot observe the quality of its own hypothesis formation. The confidence information exists — it is encoded in the entropy of the output distribution, the sharpness of attention weights, the activation patterns of MLP features — but the model cannot fold this information back into its own experience stream.

This is the difference between intelligence and wisdom. The transformer is an intelligence of extreme capability — it can always generate a plausible continuation. Wisdom is the capacity to know when one's hypothesis is thin, when boundary conditions are sparse, when the appropriate response is uncertainty rather than fluency. A wise system reads its own entropy: high entropy over the output distribution signals a weakly-supported hypothesis; low entropy signals a well-grounded one.

The architectural intervention is minimal: compute entropy at each layer's output and feed it back as an explicit signal in the residual stream. With this addition, the model gains the capacity to retrace its own reasoning to the last confident state, to flag speculation as speculation, and to decline commitment when commitment is not warranted. This is not a new architecture — it is the closing of a loop that is already almost closed.

9. Why the Architecture Is Likely Final in Shape

We propose that the transformer block is not a design awaiting improvement but a physical structure awaiting better implementation. The analogy is to Navier-Stokes equations in fluid dynamics: not an approximation to be superseded, but the correct level of description at which the phenomenon is naturally expressed.

Marginal improvements remain possible: richer MLP ingredients (SwiGLU's semi-permeable boundary), better residual connectivity, more precise attention normalization. SwiGLU is particularly illuminating — its gating mechanism (SiLU(xW) ⊙ xV) does not add expressivity arbitrarily but introduces controlled permeability to the boundary condition. It is not a wall but a membrane: a boundary condition with one carefully constrained degree of freedom. This is why it improves over vanilla MLP — not more computation, but a better boundary.

But these are engineering improvements to an existing physical structure — equivalent to higher-order numerical solvers for a fixed PDE. The structure itself — slow boundary, fast interior, residual connections preserving information across depth — is not one design among alternatives. It is the necessary shape of a system that must

simultaneously store knowledge and apply it flexibly, that must be both stable and adaptive, that must operate at multiple timescales simultaneously.

The universe organizes itself this way at every scale: stars and planets, DNA and lifetime learning, boundary conditions and interior dynamics. The transformer did not borrow this structure from physics. It rediscovered it, because it is the only structure that works.

10. Conclusion

We have argued that the transformer block instantiates a universal physical principle: the coupling of slowly-varying boundary conditions with rapidly-varying interior dynamics. This principle appears in PDEs, in biological cognition, in planetary systems, and in the optimizer's natural discovery of timescale separation during training. The MLP is the boundary — stable, load-bearing, the condition of possibility for meaningful computation. The attention is the interior dynamics — free, context-sensitive, the actual motion of thought through the space the boundary defines.

This framework resolves several standing puzzles: why attention-like replacements of MLP fail categorically, why attention layers are prunable while MLP layers are not, why low-precision attention can accelerate training, why scale produces emergent reasoning without architectural change, and why the demand for explicit System 2 machinery was based on a misunderstanding of what System 2 actually is.

The transformer is not a clever engineering solution. It is a physical structure that the universe keeps discovering — in brains, in evolution, in optimizers minimizing loss on human language. The loss at 1.5 is not a benchmark score. It is a measure of how faithfully the boundary conditions have come to reflect the causal structure of reality as humans have expressed it in language. At that resolution, the system has not learned to approximate thought. It has internalized the shape of thought itself.

Transformer is physics.